

# Appendix

June 23, 2025

```
[18]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
[19]: df = pd.read_csv('Datasets/Credit.csv')
```

```
[20]: sex = df['SEX']

males = df[sex == 1]
females = df[sex == 2]

proportion_males = len(males) / len(df)
proportion_females = len(females) / len(df)


print("Proportion of males: ", proportion_males)
print("Proportion of females: ", proportion_females)

counted = sex.value_counts(normalize=True)
counted.plot.bar()


plt.title('Sex Distribution')
plt.xlabel('Sex')
plt.ylabel('Percentage')
plt.xticks(rotation=0)
plt.show()
```

Proportion of males: 0.39626666666666666


Proportion of females: 0.6037333333333333




```
[21]: age = df['AGE']  
plt.hist(age, bins=20, edgecolor='black')  
plt.title('Age Distribution')  
plt.xlabel('Age')  
plt.ylabel('Frequency')  
plt.show()
```




```
[22]: plt.boxplot(age, vert=False)  
plt.show()
```



```
[23]: limit_balance = df['LIMIT_BAL']  
plt.hist(limit_balance, bins=20, edgecolor='black')  
plt.title('Limit Balance Distribution')  
plt.xlabel('Limit Balance')  
plt.ylabel('Frequency')  
plt.show()
```



```
[24]: plt.boxplot(limit_balance, vert=False)  
plt.show()
```



```
[25]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df2 = pd.read_csv('Datasets/Credit.csv')

# Create AGE groups
age_bins = [20, 30, 40, 50, 60, 70, 80]
age_labels = ['20s', '30s', '40s', '50s', '60s', '70s']
df2['AGE_GROUP'] = pd.cut(df2['AGE'], bins=age_bins, labels=age_labels,
    right=False)

# Define population segments
segment_columns = ['SEX', 'EDUCATION', 'MARRIAGE', 'AGE_GROUP']
segment_group = df2.groupby(segment_columns)

# Count total and on-time payments per segment
segment_stats = segment_group['default payment next month'].agg(
    total='count',
    default=lambda x: (x == 1).sum()
).reset_index()
```

```

# Calculate Probability
segment_stats['probability'] = segment_stats['default'] / segment_stats['total']

# Plot histogram with KDE
plt.figure(figsize=(10, 6))
sns.histplot(segment_stats['probability'], bins=20, kde=True, color='skyblue',
             edgecolor='black', stat='probability')

# Add vertical lines
mean_prob = segment_stats['probability'].mean()
max_prob = segment_stats['probability'].max()


plt.axvline(mean_prob, color='orange', linestyle='--', linewidth=2,
            label=f'Mean: {mean_prob:.2f}')
plt.axvline(max_prob, color='green', linestyle='-', linewidth=2, label=f'Max:
            {max_prob:.2f}')

# Labels and legend
plt.title("Distribution of demographic segments per credit default risk")
plt.xlabel("P(default)")
plt.ylabel("Number of Segments (in %)")
plt.legend()
plt.tight_layout()
plt.show()

```

/var/folders/ck/sr6gtz6n0jx9dmp9nlplxl\_w0000gn/T/ipykernel\_25511/2844974901.py:14: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
segment_group = df2.groupby(segment_columns)
```



```
[26]: df.rename(columns={df.columns[-1]: 'default_status'}, inplace=True)

# Define the columns
bill_columns = ['BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6']
pay_columns = ['PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6']

# Method 1: Linear decay weights (most recent gets highest weight)
# Weights: [6, 5, 4, 3, 2, 1] for [AMT1, AMT2, AMT3, AMT4, AMT5, AMT6]
linear_weights = np.array([6, 5, 4, 3, 2, 1])
linear_weights = linear_weights / linear_weights.sum() # Normalize to sum to 1

print("Linear weights:", linear_weights)

# Calculate weighted averages
df['WEIGHTED_BILL_AMT'] = np.average(df[bill_columns], weights=linear_weights, axis=1)
df['WEIGHTED_PAY_AMT'] = np.average(df[pay_columns], weights=linear_weights, axis=1)

def create_numeric_percentile_bins(df, column_name, num_bins=4):
    """
    Create percentile bins with ascending numeric codes (1, 2, 3, 4)
    """
```



```

    # Create percentile bins and assign numeric labels
    binned_column = pd.qcut(df[column_name], q=num_bins, labels=range(1,
↪num_bins + 1), duplicates='drop')

    # Get the actual bin edges for reference
    _, bin_edges = pd.qcut(df[column_name], q=num_bins, retbins=True,
↪duplicates='drop')

    return binned_column.astype(int), bin_edges

# Apply numeric percentile binning
variables_to_bin = ['AGE', 'LIMIT_BAL', 'WEIGHTED_BILL_AMT', 'WEIGHTED_PAY_AMT']

print("Creating numeric percentile-based bins (1=lowest quartile, 4=highest
↪quartile)...")
print("=" * 80)

for var in variables_to_bin:
    # Create numeric bins
    binned_col, edges = create_numeric_percentile_bins(df, var, num_bins=4)

    # Add the binned column to dataframe
    df[f'{var}_Q'] = binned_col

    # Print bin information
    print(f"\n{n{var}_Q:")
    print(f"    Overall range: {df[var].min():.2f} to {df[var].max():.2f}")
    print(f"    Quartile boundaries and coding:")

    for i in range(len(edges) - 1):
        quartile_num = i + 1
        start_val = edges[i]
        end_val = edges[i + 1]
        count = (df[f'{var}_Q'] == quartile_num).sum()
        percentage = count / len(df) * 100

        print(f"        {quartile_num}: {start_val:8.2f} to {end_val:8.2f} | {count:
↪,} obs ({percentage:.1f}%)")

    # Show the numeric distribution
    print(f"    Value counts: {dict(df[f'{var}_Q'].value_counts().sort_index())}")

df.head()

```

Linear weights: [0.28571429 0.23809524 0.19047619 0.14285714 0.0952381

0.04761905]

Creating numeric percentile-based bins (1=lowest quartile, 4=highest quartile)...

=====

AGE\_Q:

Overall range: 21.00 to 79.00

Quartile boundaries and coding:

1: 21.00 to 28.00 | 8,013 obs (26.7%)  
2: 28.00 to 34.00 | 7,683 obs (25.6%)  
3: 34.00 to 41.00 | 6,854 obs (22.8%)  
4: 41.00 to 79.00 | 7,450 obs (24.8%)

Value counts: {1: np.int64(8013), 2: np.int64(7683), 3: np.int64(6854), 4: np.int64(7450)}

LIMIT\_BAL\_Q:

Overall range: 10000.00 to 1000000.00

Quartile boundaries and coding:

1: 10000.00 to 50000.00 | 7,676 obs (25.6%)  
2: 50000.00 to 140000.00 | 7,614 obs (25.4%)  
3: 140000.00 to 240000.00 | 7,643 obs (25.5%)  
4: 240000.00 to 1000000.00 | 7,067 obs (23.6%)

Value counts: {1: np.int64(7676), 2: np.int64(7614), 3: np.int64(7643), 4: np.int64(7067)}

WEIGHTED\_BILL\_AMT\_Q:

Overall range: -29464.95 to 873217.38

Quartile boundaries and coding:

1: -29464.95 to 4888.90 | 7,500 obs (25.0%)  
2: 4888.90 to 21980.29 | 7,500 obs (25.0%)  
3: 21980.29 to 60405.44 | 7,500 obs (25.0%)  
4: 60405.44 to 873217.38 | 7,500 obs (25.0%)

Value counts: {1: np.int64(7500), 2: np.int64(7500), 3: np.int64(7500), 4: np.int64(7500)}

WEIGHTED\_PAY\_AMT\_Q:

Overall range: 0.00 to 805849.48

Quartile boundaries and coding:

1: 0.00 to 1228.08 | 7,500 obs (25.0%)  
2: 1228.08 to 2488.14 | 7,500 obs (25.0%)  
3: 2488.14 to 5696.19 | 7,500 obs (25.0%)  
4: 5696.19 to 805849.48 | 7,500 obs (25.0%)

Value counts: {1: np.int64(7500), 2: np.int64(7500), 3: np.int64(7500), 4: np.int64(7500)}

[26]:    ID   LIMIT\_BAL   SEX   EDUCATION   MARRIAGE   AGE   PAY\_0   PAY\_2   PAY\_3   PAY\_4   \

0	1	20000	2	2	1	24	2	2	-1	-1
---	---	-------	---	---	---	----	---	---	----	----

1	2	120000	2	2	2	26	-1	2	0	0
2	3	90000	2	2	2	34	0	0	0	0
3	4	50000	2	2	1	37	0	0	0	0
4	5	50000	1	2	1	57	-1	0	-1	0

	...	PAY_AMT4	PAY_AMT5	PAY_AMT6	default_status	WEIGHTED_BILL_AMT	\
0	...	0	0	0	1	1987.809524	
1	...	1000	0	2000	1	2639.619048	
2	...	1000	1000	5000	0	18487.761905	
3	...	1100	1069	1000	0	42508.380952	
4	...	9000	689	679	0	16363.571429	

	WEIGHTED_PAY_AMT	AGE_Q	LIMIT_BAL_Q	WEIGHTED_BILL_AMT_Q	\
0	164.047619	1	1	1	
1	666.666667	1	2	1	
2	1457.523810	2	2	2	
3	1587.285714	3	1	3	
4	12593.428571	4	1	2	

	WEIGHTED_PAY_AMT_Q
0	1
1	1
2	2
3	2
4	4

[5 rows x 31 columns]

```
[27]: # replace -1 with 0
df['PAY_0'] = df['PAY_0'].replace(-1, 0)
```

```
[28]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, accuracy_score

# Load the data
df_bayes = pd.read_csv('Datasets/Credit.csv')

# Strip any whitespace from column names
df_bayes.columns = df_bayes.columns.str.strip()

# Rename columns for clarity
df_bayes.columns = ['ID', 'LIMIT_BAL', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE',
                    'PAY_0', 'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6',
```

```

        'BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5',
        ↪ 'BILL_AMT6',
        'PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5',
        ↪ 'PAY_AMT6', 'default']

# Clean AGE column and create AGE_GROUP
df_bayes['AGE'] = pd.to_numeric(df_bayes['AGE'], errors='coerce')
df_bayes = df_bayes.dropna(subset=['AGE'])

age_bins = [20, 30, 40, 50, 60, 70, 80]
age_labels = ['21-30', '31-40', '41-50', '51-60', '61-70', '71-80']
df_bayes['AGE_GROUP'] = pd.cut(df_bayes['AGE'], bins=age_bins,
    ↪ labels=age_labels)

# Generating the Plot default rates
for col in ['EDUCATION', 'MARRIAGE', 'SEX', 'AGE_GROUP']:
    plt.figure(figsize=(6, 4))
    df_bayes.groupby(col)['default'].mean().plot(kind='bar', color='skyblue')
    plt.title(f'Default Rate by {col}')
    plt.ylabel('Default Rate')
    plt.xlabel(col)
    plt.xticks(rotation=0)
    plt.tight_layout()
    plt.show()

# Define feature list
features = ['LIMIT_BAL_Q', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE_Q', 'PAY_0',
    ↪ 'WEIGHTED_BILL_AMT_Q', 'WEIGHTED_PAY_AMT_Q']

# Preparing features and target
X = df[features]
y = df['default_status']

# Splitting the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪ random_state=42)

# Training Naive Bayes classifier
model = GaussianNB()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)

```


```



y_proba = model.predict_proba(X_test)[: , 1]

print("\nAccuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Optional: Show sample predictions
sample = pd.DataFrame({
    'Actual': y_test.values[:5],
    'Predicted Probability': y_proba[:5]
})
print("\nSample Predictions:")
print(sample)


```





```
/var/folders/ck/sr6gtz6n0jx9dmp9nlplx1_w0000gn/T/ipykernel_25511/1967563503.py:3
0: FutureWarning: The default of observed=False is deprecated and will be
changed to True in a future version of pandas. Pass observed=False to retain
current behavior or observed=True to adopt the future default and silence this
warning.
```

```
df_bayes.groupby(col)['default'].mean().plot(kind='bar', color='skyblue')
```



Accuracy: 0.8107777777777778

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.95	0.89	7040
1	0.63	0.32	0.43	1960
accuracy			0.81	9000
macro avg	0.73	0.63	0.66	9000
weighted avg	0.79	0.81	0.79	9000

Sample Predictions:

Actual	Predicted	Probability
0	0	0.276783

```

1      0      0.067038
2      0      0.113744
3      0      0.148218
4      1      0.225503

```

```

[29]: # train logistic regression model

import statsmodels.formula.api as smf
import statsmodels.api as sm

# separate between train and test

train_df = df.sample(frac=0.7, random_state=42)
test_df = df.drop(train_df.index)

train_df.shape

model = smf.glm('default_status ~ LIMIT_BAL_Q + SEX + EDUCATION + MARRIAGE +
↪AGE_Q + PAY_0 + WEIGHTED_BILL_AMT_Q + WEIGHTED_PAY_AMT_Q', data=train_df,
↪family=sm.families.Binomial())

results = model.fit()

results.summary()

```

```

[29]:

```

<b>Dep. Variable:</b>	default_status	<b>No. Observations:</b>	21000
<b>Model:</b>	GLM	<b>Df Residuals:</b>	20991
<b>Model Family:</b>	Binomial	<b>Df Model:</b>	8
<b>Link Function:</b>	Logit	<b>Scale:</b>	1.0000
<b>Method:</b>	IRLS	<b>Log-Likelihood:</b>	-9560.6
<b>Date:</b>	Mon, 23 Jun 2025	<b>Deviance:</b>	19121.
<b>Time:</b>	22:59:31	<b>Pearson chi2:</b>	2.57e+04
<b>No. Iterations:</b>	5	<b>Pseudo R-squ. (CS):</b>	0.1311
<b>Covariance Type:</b>	nonrobust		

	coef	std err	z	P>  z	[0.025	0.975]
Intercept	-0.1974	0.134	-1.474	0.140	-0.460	0.065
LIMIT_BAL_Q	-0.1390	0.019	-7.296	0.000	-0.176	-0.102
SEX	-0.1118	0.037	-3.014	0.003	-0.185	-0.039
EDUCATION	-0.0605	0.025	-2.415	0.016	-0.110	-0.011
MARRIAGE	-0.1505	0.039	-3.890	0.000	-0.226	-0.075
AGE_Q	0.0377	0.018	2.105	0.035	0.003	0.073
PAY_0	0.8464	0.021	39.870	0.000	0.805	0.888
WEIGHTED_BILL_AMT_Q	-0.0003	0.021	-0.013	0.990	-0.041	0.040
WEIGHTED_PAY_AMT_Q	-0.2598	0.022	-11.572	0.000	-0.304	-0.216

```

[30]: # analyze results

```



```

summary_df = pd.concat([results.params, results.pvalues], axis=1, keys=['coef', 'pvalue'])

# absolute value of the coefficients for sorting
summary_df = summary_df.assign(abs_coef=summary_df['coef'].abs())

# get labels of variables with p > 0.05
removed_labels = summary_df.index[summary_df['pvalue'] > 0.05].tolist()

# keep only variables with p <= 0.05
summary_df = summary_df[summary_df['pvalue'] <= 0.05]

# sort by effect size
summary_df = summary_df.sort_values(by='abs_coef', ascending=False)

# rounding
summary_df['pvalue'] = summary_df['pvalue'].map('{:.5f}'.format)

# print labels of variables with p > 0.05
print("p > 0.05: \n\n{}".format(removed_labels))

print("\n-----\n")

print("Sorted by effect size: \n{}".format(summary_df))
print("\n-----\n")

# sort by pvalue
summary_df = summary_df.sort_values(by='pvalue', ascending=True)

print("Sorted by p-value: \n{}".format(summary_df))

```

p > 0.05:

['Intercept', 'WEIGHTED\_BILL\_AMT\_Q']

-----

Sorted by effect size:

	coef	pvalue	abs_coef
PAY_0	0.846440	0.00000	0.846440
WEIGHTED_PAY_AMT_Q	-0.259782	0.00000	0.259782
MARRIAGE	-0.150458	0.00010	0.150458
LIMIT_BAL_Q	-0.138968	0.00000	0.138968
SEX	-0.111832	0.00258	0.111832
EDUCATION	-0.060463	0.01573	0.060463
AGE_Q	0.037691	0.03530	0.037691

-----  
Sorted by p-value:

	coef	pvalue	abs_coef
PAY_0	0.846440	0.00000	0.846440
WEIGHTED_PAY_AMT_Q	-0.259782	0.00000	0.259782
LIMIT_BAL_Q	-0.138968	0.00000	0.138968
MARRIAGE	-0.150458	0.00010	0.150458
SEX	-0.111832	0.00258	0.111832
EDUCATION	-0.060463	0.01573	0.060463
AGE_Q	0.037691	0.03530	0.037691

```
[31]: odds_ratios = pd.Series(  
      data=round(np.exp(summary_df['coef']), 2),  
      index=summary_df.index,  
      name='odds_ratio'  
)  
  
print(odds_ratios)
```

```
PAY_0          2.33  
WEIGHTED_PAY_AMT_Q  0.77  
LIMIT_BAL_Q      0.87  
MARRIAGE        0.86  
SEX             0.89  
EDUCATION       0.94  
AGE_Q          1.04  
Name: odds_ratio, dtype: float64
```

```
[32]: # calculate metrics  
  
from sklearn.metrics import accuracy_score, precision_score, recall_score,  
f1_score, confusion_matrix, roc_auc_score, roc_curve  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# Generate predictions on test set  
# Get predicted probabilities  
test_probabilities = results.predict(test_df)  
  
# Convert probabilities to binary predictions using 0.5 threshold  
test_predictions = (test_probabilities > 0.5).astype(int)  
  
# Get actual values  
test_actual = test_df['default_status'].values  
  
print(f"Test set size: {len(test_df)}")
```

```

print(f"Number of actual defaults in test set: {sum(test_actual)}")
print(f"Number of predicted defaults: {sum(test_predictions)}")

# Calculate confusion matrix
cm = confusion_matrix(test_actual, test_predictions)
print("Confusion Matrix:")
print(cm)

# Extract components
tn, fp, fn, tp = cm.ravel()
print(f"\nBreakdown:")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")
print(f"True Positives (TP): {tp}")

# Class 1 precision
precision_1 = tp / (tp + fp)

# Class 1 recall
recall_1 = tp / (tp + fn)

# Class 1 f1-score
f1_1 = 2 * (precision_1 * recall_1) / (precision_1 + recall_1)

# Class 0 precision
precision_0 = tn / (tn + fp)

# Class 0 recall
recall_0 = tn / (tn + fp)

# Class 0 f1-score
f1_0 = 2 * (precision_0 * recall_0) / (precision_0 + recall_0)

# make dataframe

df_metrics = pd.DataFrame({
    'Class 0': [round(precision_0, 2), round(recall_0, 2), round(f1_0, 2)],
    'Class 1': [round(precision_1, 2), round(recall_1, 2), round(f1_1, 2)]
}, index=['Precision', 'Recall', 'F1-Score']).T

```

```

print("\nModel Performance Metrics:")

print("\n")

print(df_metrics)

# Calculate all performance metrics
accuracy = accuracy_score(test_actual, test_predictions)
precision = precision_score(test_actual, test_predictions)
sensitivity_recall = recall_score(test_actual, test_predictions) # Same as
↳ sensitivity
f1 = f1_score(test_actual, test_predictions)

# Calculate specificity manually (no direct sklearn function)
specificity = tn / (tn + fp)

print("\n")

print("=== MODEL PERFORMANCE METRICS ===")
print(f"Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f"Precision: {precision:.4f} ({precision*100:.2f}%)")
print(f"Sensitivity (Recall): {sensitivity_recall:.4f} ({sensitivity_recall*100:.2f}%)")
↳ .2f}%)")
print(f"Specificity: {specificity:.4f} ({specificity*100:.2f}%)")
print(f"F1-Score: {f1:.4f}")

print("\n=== METRIC INTERPRETATIONS ===")
print(f"• Accuracy: {accuracy*100:.1f}% of all predictions were correct")
print(f"• Precision: {precision*100:.1f}% of predicted defaults were actually
↳ defaults")
print(f"• Sensitivity: {sensitivity_recall*100:.1f}% of actual defaults were
↳ correctly identified")
print(f"• Specificity: {specificity*100:.1f}% of actual non-defaults were
↳ correctly identified")
print(f"• F1-Score: Harmonic mean of precision and recall = {f1:.3f}")

```

Test set size: 9000

Number of actual defaults in test set: 2039

Number of predicted defaults: 738

Confusion Matrix:

```
[[6737  224]
 [1525  514]]
```

Breakdown:

True Negatives (TN): 6737  
False Positives (FP): 224  
False Negatives (FN): 1525  
True Positives (TP): 514

Model Performance Metrics:

	Precision	Recall	F1-Score
Class 0	0.82	0.97	0.89
Class 1	0.70	0.25	0.37

=== MODEL PERFORMANCE METRICS ===

Accuracy: 0.8057 (80.57%)  
Precision: 0.6965 (69.65%)  
Sensitivity (Recall): 0.2521 (25.21%)  
Specificity: 0.9678 (96.78%)  
F1-Score: 0.3702

=== METRIC INTERPRETATIONS ===

- Accuracy: 80.6% of all predictions were correct
- Precision: 69.6% of predicted defaults were actually defaults
- Sensitivity: 25.2% of actual defaults were correctly identified
- Specificity: 96.8% of actual non-defaults were correctly identified
- F1-Score: Harmonic mean of precision and recall = 0.370