

FinalTeamProject

June 21, 2025

```
[216]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
[217]: df = pd.read_csv('Datasets/Credit.csv')
```


```
[218]: # Load the dataset -- Skip first metadata row
df = pd.read_csv('Datasets/Credit.csv')

# Drop the ID column as it's not useful for analysis
df_cleaned = df.drop(columns=["ID"])

# Set Seaborn style
sns.set(style="whitegrid")

# Generating Correlation Heatmap
plt.figure(figsize=(18, 14))
correlation_matrix = df_cleaned.corr()
sns.heatmap(correlation_matrix, annot=False, cmap="coolwarm", fmt=".2f",
            linewidths=0.5)
plt.title("Correlation Heatmap")
plt.tight_layout()
plt.show()

# Printing my explanation
print("explanation")
print()
print("The correlation heatmap of the dataset reveals relationships between
    ↳ features such as: High correlation among BILL_AMT variables (e.g.,
    ↳ BILL_AMT1, BILL_AMT2, etc.)")
print("The correlation heatmap also indicates a positive correlation between
    ↳ LIMIT_BAL and PAY_AMT values, a Mmoderate positive correlation between past
    ↳ payment statuses (PAY_0 to PAY_6) and default likelihood")
```



explanation

The correlation heatmap of the dataset reveals relationships between features such as: High correlation among BILL_AMT variables (e.g., BILL_AMT1, BILL_AMT2, etc.)

The correlation heatmap also indicates a positive correlation between LIMIT_BAL and PAY_AMT values, a Moderate positive correlation between past payment statuses (PAY_0 to PAY_6) and default likelihood

```
[219]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset (assuming the file is in the same directory) --Skip first_
↳ metadata row
df = pd.read_csv('Datasets/Credit.csv')

# Drop the ID column as it's not useful for analysis
```

```


df_cleaned = df.drop(columns=["ID"])

# Set Seaborn style
sns.set(style="whitegrid")

# --- Histogram of Credit Limit ---
plt.figure(figsize=(10, 6))
sns.histplot(df_cleaned["LIMIT_BAL"], bins=30, kde=True, color="skyblue")
plt.title("Distribution of Credit Limit (LIMIT_BAL)")
plt.xlabel("Credit Limit")
plt.ylabel("Frequency")
plt.tight_layout()
plt.show()

# Printing my explanation
print("Explanation -- The histogram of the credit limit (LIMIT_BAL)")
print()
print("Most credit limits are concentrated below 200,000 units.")
print("The distribution is right-skewed, indicating a smaller number of clients_
↳with very high credit limits.")

```



Explanation -- The histogram of the credit limit (LIMIT_BAL)

Most credit limits are concentrated below 200,000 units.

The distribution is right-skewed, indicating a smaller number of clients with very high credit limits.

```

[220]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df2 = pd.read_csv('Datasets/Credit.csv')

# Create AGE groups
age_bins = [20, 30, 40, 50, 60, 70, 80]
age_labels = ['20s', '30s', '40s', '50s', '60s', '70s']
df2['AGE_GROUP'] = pd.cut(df2['AGE'], bins=age_bins, labels=age_labels,
    ↪right=False)

# Define population segments
segment_columns = ['SEX', 'EDUCATION', 'MARRIAGE', 'AGE_GROUP']
segment_group = df2.groupby(segment_columns)

# Count total and on-time payments per segment
segment_stats = segment_group['default payment next month'].agg(
    total='count',
    default=lambda x: (x == 1).sum()
).reset_index()

# Calculate Probability
segment_stats['probability'] = segment_stats['default'] / segment_stats['total']

# Plot histogram with KDE
plt.figure(figsize=(10, 6))
sns.histplot(segment_stats['probability'], bins=20, kde=True, color='skyblue',
    ↪edgecolor='black', stat='probability')

# Add vertical lines
mean_prob = segment_stats['probability'].mean()
max_prob = segment_stats['probability'].max()

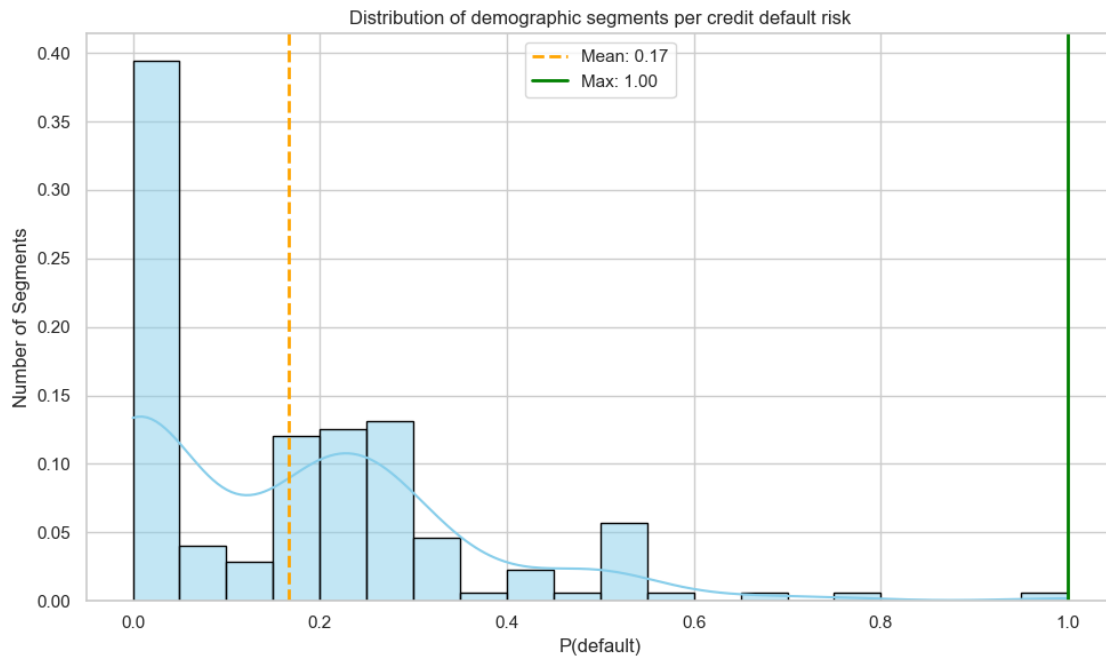
plt.axvline(mean_prob, color='orange', linestyle='--', linewidth=2,
    ↪label=f'Mean: {mean_prob:.2f}')
plt.axvline(max_prob, color='green', linestyle='-', linewidth=2, label=f'Max:
    ↪{max_prob:.2f}')

# Labels and legend
plt.title("Distribution of demographic segments per credit default risk")
plt.xlabel("P(default)")
plt.ylabel("Number of Segments")
plt.legend()
plt.tight_layout()
plt.show()

```

```
/var/folders/ck/sr6gtz6n0jx9dmp9nlplxl_w0000gn/T/ipykernel_65554/2297665363.py:1
4: FutureWarning: The default of observed=False is deprecated and will be
changed to True in a future version of pandas. Pass observed=False to retain
current behavior or observed=True to adopt the future default and silence this
warning.
```

```
segment_group = df2.groupby(segment_columns)
```



```
[221]: df.rename(columns={df.columns[-1]: 'default_status'}, inplace=True)

# Define the columns
bill_columns = ['BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6']
pay_columns = ['PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6']

# Method 1: Linear decay weights (most recent gets highest weight)
# Weights: [6, 5, 4, 3, 2, 1] for [AMT1, AMT2, AMT3, AMT4, AMT5, AMT6]
linear_weights = np.array([6, 5, 4, 3, 2, 1])
linear_weights = linear_weights / linear_weights.sum() # Normalize to sum to 1

print("Linear weights:", linear_weights)

# Calculate weighted averages
df['WEIGHTED_BILL_AMT'] = np.average(df[bill_columns], weights=linear_weights, axis=1)
```

```

df['WEIGHTED_PAY_AMT'] = np.average(df[pay_columns], weights=linear_weights,
    ↪axis=1)

def create_numeric_percentile_bins(df, column_name, num_bins=4):
    """
    Create percentile bins with ascending numeric codes (1, 2, 3, 4)
    """
    # Create percentile bins and assign numeric labels
    binned_column = pd.qcut(df[column_name], q=num_bins, labels=range(1,
    ↪num_bins + 1), duplicates='drop')

    # Get the actual bin edges for reference
    _, bin_edges = pd.qcut(df[column_name], q=num_bins, retbins=True,
    ↪duplicates='drop')

    return binned_column.astype(int), bin_edges

# Apply numeric percentile binning
variables_to_bin = ['AGE', 'LIMIT_BAL', 'WEIGHTED_BILL_AMT', 'WEIGHTED_PAY_AMT']

print("Creating numeric percentile-based bins (1=lowest quartile, 4=highest
    ↪quartile)...")
print("=" * 80)

for var in variables_to_bin:
    # Create numeric bins
    binned_col, edges = create_numeric_percentile_bins(df, var, num_bins=4)

    # Add the binned column to dataframe
    df[f'{var}_Q'] = binned_col

    # Print bin information
    print(f"\n{var}_Q:")
    print(f" Overall range: {df[var].min():.2f} to {df[var].max():.2f}")
    print(f" Quartile boundaries and coding:")

    for i in range(len(edges) - 1):
        quartile_num = i + 1
        start_val = edges[i]
        end_val = edges[i + 1]
        count = (df[f'{var}_Q'] == quartile_num).sum()
        percentage = count / len(df) * 100

        print(f" {quartile_num}: {start_val:8.2f} to {end_val:8.2f} | {count:
    ↪,} obs ({percentage:.1f}%)")

    # Show the numeric distribution

```

```

print(f" Value counts: {dict(df[f'{var}_Q'].value_counts().sort_index())}")

df.head()

```

Linear weights: [0.28571429 0.23809524 0.19047619 0.14285714 0.0952381
0.04761905]

Creating numeric percentile-based bins (1=lowest quartile, 4=highest
quartile)...

AGE_Q:

Overall range: 21.00 to 79.00

Quartile boundaries and coding:

1:	21.00 to	28.00	8,013 obs (26.7%)
2:	28.00 to	34.00	7,683 obs (25.6%)
3:	34.00 to	41.00	6,854 obs (22.8%)
4:	41.00 to	79.00	7,450 obs (24.8%)

Value counts: {1: np.int64(8013), 2: np.int64(7683), 3: np.int64(6854), 4:
np.int64(7450)}

LIMIT_BAL_Q:

Overall range: 10000.00 to 1000000.00

Quartile boundaries and coding:

1:	10000.00 to	50000.00	7,676 obs (25.6%)
2:	50000.00 to	140000.00	7,614 obs (25.4%)
3:	140000.00 to	240000.00	7,643 obs (25.5%)
4:	240000.00 to	1000000.00	7,067 obs (23.6%)

Value counts: {1: np.int64(7676), 2: np.int64(7614), 3: np.int64(7643), 4:
np.int64(7067)}

WEIGHTED_BILL_AMT_Q:

Overall range: -29464.95 to 873217.38

Quartile boundaries and coding:

1:	-29464.95 to	4888.90	7,500 obs (25.0%)
2:	4888.90 to	21980.29	7,500 obs (25.0%)
3:	21980.29 to	60405.44	7,500 obs (25.0%)
4:	60405.44 to	873217.38	7,500 obs (25.0%)

Value counts: {1: np.int64(7500), 2: np.int64(7500), 3: np.int64(7500), 4:
np.int64(7500)}

WEIGHTED_PAY_AMT_Q:

Overall range: 0.00 to 805849.48

Quartile boundaries and coding:

1:	0.00 to	1228.08	7,500 obs (25.0%)
2:	1228.08 to	2488.14	7,500 obs (25.0%)
3:	2488.14 to	5696.19	7,500 obs (25.0%)

4: 5696.19 to 805849.48 | 7,500 obs (25.0%)

Value counts: {1: np.int64(7500), 2: np.int64(7500), 3: np.int64(7500), 4: np.int64(7500)}

```
[221]:
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	\
0	1	20000	2	2	1	24	2	2	-1	-1	
1	2	120000	2	2	2	26	-1	2	0	0	
2	3	90000	2	2	2	34	0	0	0	0	
3	4	50000	2	2	1	37	0	0	0	0	
4	5	50000	1	2	1	57	-1	0	-1	0	

	...	PAY_AMT4	PAY_AMT5	PAY_AMT6	default_status	WEIGHTED_BILL_AMT	\
0	...	0	0	0	1	1987.809524	
1	...	1000	0	2000	1	2639.619048	
2	...	1000	1000	5000	0	18487.761905	
3	...	1100	1069	1000	0	42508.380952	
4	...	9000	689	679	0	16363.571429	

	WEIGHTED_PAY_AMT	AGE_Q	LIMIT_BAL_Q	WEIGHTED_BILL_AMT_Q	\
0	164.047619	1	1	1	
1	666.666667	1	2	1	
2	1457.523810	2	2	2	
3	1587.285714	3	1	3	
4	12593.428571	4	1	2	

	WEIGHTED_PAY_AMT_Q
0	1
1	1
2	2
3	2
4	4

[5 rows x 31 columns]

```
[222]: # replace -1 with 0
df['PAY_0'] = df['PAY_0'].replace(-1, 0)

# separate between train and test

train_df = df.sample(frac=0.7, random_state=42)
test_df = df.drop(train_df.index)

train_df.shape
```

[222]: (21000, 31)


```
[223]: # train logistic regression model

import statsmodels.formula.api as smf
import statsmodels.api as sm

model = smf.glm('default_status ~ LIMIT_BAL_Q + SEX + EDUCATION + MARRIAGE +
↪AGE_Q + PAY_0 + WEIGHTED_BILL_AMT_Q + WEIGHTED_PAY_AMT_Q', data=train_df,
↪family=sm.families.Binomial())

results = model.fit()

results.summary()
```

[223]:

Dep. Variable:	default_status	No. Observations:	21000
Model:	GLM	Df Residuals:	20991
Model Family:	Binomial	Df Model:	8
Link Function:	Logit	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-9560.6
Date:	Sat, 21 Jun 2025	Deviance:	19121.
Time:	13:57:33	Pearson chi2:	2.57e+04
No. Iterations:	5	Pseudo R-squ. (CS):	0.1311
Covariance Type:	nonrobust		

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1974	0.134	-1.474	0.140	-0.460	0.065
LIMIT_BAL_Q	-0.1390	0.019	-7.296	0.000	-0.176	-0.102
SEX	-0.1118	0.037	-3.014	0.003	-0.185	-0.039
EDUCATION	-0.0605	0.025	-2.415	0.016	-0.110	-0.011
MARRIAGE	-0.1505	0.039	-3.890	0.000	-0.226	-0.075
AGE_Q	0.0377	0.018	2.105	0.035	0.003	0.073
PAY_0	0.8464	0.021	39.870	0.000	0.805	0.888
WEIGHTED_BILL_AMT_Q	-0.0003	0.021	-0.013	0.990	-0.041	0.040
WEIGHTED_PAY_AMT_Q	-0.2598	0.022	-11.572	0.000	-0.304	-0.216

```
[224]: # analyze results

summary_df = pd.concat([results.params, results.pvalues], axis=1, keys=['coef',
↪'pvalue'])

# absolute value of the coefficients for sorting
summary_df = summary_df.assign(abs_coef=summary_df['coef'].abs())

# get labels of variables with p > 0.05
removed_labels = summary_df.index[summary_df['pvalue'] > 0.05].tolist()
```

```

# keep only variables with p <= 0.05
summary_df = summary_df[summary_df['pvalue'] <= 0.05]

# sort by effect size
summary_df = summary_df.sort_values(by='abs_coef', ascending=False)

# rounding
summary_df['pvalue'] = summary_df['pvalue'].map('{:.5f}'.format)

# print labels of variables with p > 0.05
print("p > 0.05: \n\n{}".format(removed_labels))

print("\n-----\n")

print("Sorted by effect size: \n{}".format(summary_df))
print("\n-----\n")

# sort by pvalue
summary_df = summary_df.sort_values(by='pvalue', ascending=True)

print("\n-----\n")

print("Sorted by p-value: \n{}".format(summary_df))
print("\n-----\n")

```

p > 0.05:

['Intercept', 'WEIGHTED_BILL_AMT_Q']

Sorted by effect size:

	coef	pvalue	abs_coef
PAY_0	0.846440	0.00000	0.846440
WEIGHTED_PAY_AMT_Q	-0.259782	0.00000	0.259782
MARRIAGE	-0.150458	0.00010	0.150458
LIMIT_BAL_Q	-0.138968	0.00000	0.138968
SEX	-0.111832	0.00258	0.111832
EDUCATION	-0.060463	0.01573	0.060463
AGE_Q	0.037691	0.03530	0.037691

Sorted by p-value:

	coef	pvalue	abs_coef
--	------	--------	----------

PAY_0	0.846440	0.00000	0.846440
WEIGHTED_PAY_AMT_Q	-0.259782	0.00000	0.259782
LIMIT_BAL_Q	-0.138968	0.00000	0.138968
MARRIAGE	-0.150458	0.00010	0.150458
SEX	-0.111832	0.00258	0.111832
EDUCATION	-0.060463	0.01573	0.060463
AGE_Q	0.037691	0.03530	0.037691

```
[225]: odds_ratios = pd.Series(
        data=round(np.exp(summary_df['coef']), 2),
        index=summary_df.index,
        name='odds_ratio'
    )

    print(odds_ratios)
```

PAY_0	2.33
WEIGHTED_PAY_AMT_Q	0.77
LIMIT_BAL_Q	0.87
MARRIAGE	0.86
SEX	0.89
EDUCATION	0.94
AGE_Q	1.04

Name: odds_ratio, dtype: float64

```
[226]: # Make examples

class Person:

    def __init__(self, age, sex, education, marriage, limit_balance,
↪bill_amount, payment_amount, payment_history):
        self.age = age
        self.sex = sex
        self.education = education
        self.marriage = marriage
        self.limit_balance = limit_balance
        self.bill_amount = bill_amount
        self.payment_amount = payment_amount
        self.payment_history = payment_history

    def calculate_probability(self):
        intercept = results.params['Intercept']
        age_coef = results.params['AGE_Q']
        sex_coef = results.params['SEX']
        education_coef = results.params['EDUCATION']
```

```

marriage_coef = results.params['MARRIAGE']
limit_balance_coef = results.params['LIMIT_BAL_Q']
bill_amount_coef = results.params['WEIGHTED_BILL_AMT_Q']
payment_amount_coef = results.params['WEIGHTED_PAY_AMT_Q']
payment_history_coef = results.params['PAY_0']

probability = 1 / (1 + np.exp(-(intercept + age_coef * self.age +
↪sex_coef * self.sex + education_coef * self.education + marriage_coef * self.
↪marriage + limit_balance_coef * self.limit_balance + bill_amount_coef * self.
↪bill_amount + payment_amount_coef * self.payment_amount +
↪payment_history_coef * self.payment_history)))

return probability

jake = Person(age=1, sex=1, education=0, marriage=0, limit_balance=1,
↪bill_amount=2, payment_amount=0, payment_history=0)
print("jake:", round(jake.calculate_probability(), 4))

john = Person(age=1, sex=1, education=4, marriage=3, limit_balance=1,
↪bill_amount=4, payment_amount=0, payment_history=8)
print("john:", round(john.calculate_probability(), 4))

penelope = Person(age=4, sex=2, education=1, marriage=1, limit_balance=4,
↪bill_amount=1, payment_amount=3, payment_history=0)
print("penelope:", round(penelope.calculate_probability(), 4))

ricardo = Person(age=1, sex=1, education=1, marriage=0, limit_balance=4,
↪bill_amount=4, payment_amount=1, payment_history=6)
print("ricardo:", round(ricardo.calculate_probability(), 4))

stella = Person(age=2, sex=2, education=3, marriage=2, limit_balance=1,
↪bill_amount=1, payment_amount=1, payment_history=0)
print("stella:", round(stella.calculate_probability(), 4))

```

```

jake: 0.3987
john: 0.9966
penelope: 0.1398
ricardo: 0.9807
stella: 0.2267

```

```

[227]: # calculate metrics

from sklearn.metrics import accuracy_score, precision_score, recall_score,
↪f1_score, confusion_matrix, roc_auc_score, roc_curve

```

```

import matplotlib.pyplot as plt
import seaborn as sns

# Generate predictions on test set
# Get predicted probabilities
test_probabilities = results.predict(test_df)

# Convert probabilities to binary predictions using 0.5 threshold
test_predictions = (test_probabilities > 0.5).astype(int)

# Get actual values
test_actual = test_df['default_status'].values

print(f"Test set size: {len(test_df)}")
print(f"Number of actual defaults in test set: {sum(test_actual)}")
print(f"Number of predicted defaults: {sum(test_predictions)}")

# Calculate confusion matrix
cm = confusion_matrix(test_actual, test_predictions)
print("Confusion Matrix:")
print(cm)

# Extract components
tn, fp, fn, tp = cm.ravel()
print(f"\nBreakdown:")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")
print(f"True Positives (TP): {tp}")

# Calculate all performance metrics
accuracy = accuracy_score(test_actual, test_predictions)
precision = precision_score(test_actual, test_predictions)
sensitivity_recall = recall_score(test_actual, test_predictions) # Same as sensitivity
f1 = f1_score(test_actual, test_predictions)

# Calculate specificity manually (no direct sklearn function)
specificity = tn / (tn + fp)

print("=== MODEL PERFORMANCE METRICS ===")
print(f"Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f"Precision: {precision:.4f} ({precision*100:.2f}%)")
print(f"Sensitivity (Recall): {sensitivity_recall:.4f} ({sensitivity_recall*100:.2f}%)")

```

```

print(f"Specificity: {specificity:.4f} ({specificity*100:.2f}%)")
print(f"F1-Score: {f1:.4f}")

print("\n=== METRIC INTERPRETATIONS ===")
print(f"• Accuracy: {accuracy*100:.1f}% of all predictions were correct")
print(f"• Precision: {precision*100:.1f}% of predicted defaults were actually defaults")
print(f"• Sensitivity: {sensitivity_recall*100:.1f}% of actual defaults were correctly identified")
print(f"• Specificity: {specificity*100:.1f}% of actual non-defaults were correctly identified")
print(f"• F1-Score: Harmonic mean of precision and recall = {f1:.3f}")

# Calculate AUC
auc = roc_auc_score(test_actual, test_probabilities)
print(f"AUC-ROC Score: {auc:.4f}")

# Generate ROC curve data
fpr, tpr, thresholds = roc_curve(test_actual, test_probabilities)

# Plot ROC curve
plt.figure(figsize=(10, 8))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC Curve (AUC = {auc:.3f})')
plt.plot([0, 1], [0, 1], color='red', lw=2, linestyle='--', label='Random Classifier (AUC = 0.5)')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.title('ROC Curve - Credit Default Prediction Model')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

print(f"\n=== AUC INTERPRETATION ===")
if auc >= 0.9:
    interpretation = "Excellent"
elif auc >= 0.8:
    interpretation = "Good"
elif auc >= 0.7:
    interpretation = "Fair"
elif auc >= 0.6:
    interpretation = "Poor"
else:
    interpretation = "Very Poor"

print(f"AUC = {auc:.3f} indicates {interpretation} discriminatory ability")

```

Test set size: 9000
Number of actual defaults in test set: 2039
Number of predicted defaults: 738
Confusion Matrix:
[[6737 224]
 [1525 514]]


Breakdown:

True Negatives (TN): 6737
False Positives (FP): 224
False Negatives (FN): 1525
True Positives (TP): 514
=== MODEL PERFORMANCE METRICS ===
Accuracy: 0.8057 (80.57%)
Precision: 0.6965 (69.65%)
Sensitivity (Recall): 0.2521 (25.21%)
Specificity: 0.9678 (96.78%)
F1-Score: 0.3702

=== METRIC INTERPRETATIONS ===

- Accuracy: 80.6% of all predictions were correct
- Precision: 69.6% of predicted defaults were actually defaults
- Sensitivity: 25.2% of actual defaults were correctly identified
- Specificity: 96.8% of actual non-defaults were correctly identified
- F1-Score: Harmonic mean of precision and recall = 0.370

AUC-ROC Score: 0.7326



=== AUC INTERPRETATION ===

AUC = 0.733 indicates Fair discriminatory ability